

Algorithms and Distributed Systems 2023/2024 (Lecture One)

**MIEI - Integrated Master in Computer Science and
Informatics**

**MEI – Master in Computer Science and
Informatics**

Specialization block

Nuno Preguiça (nmp@fct.unl.pt)

Alex Davidson (a.davidson@fct.unl.pt)



NOVA SCHOOL OF
SCIENCE & TECHNOLOGY

Based on slides from João Leitão

Lecture structure:

- Understand what is a distributed system.
- Modeling a distributed system.
 - General system model.
 - Fault model.
 - Timing assumption.
- Understanding the role of an algorithm.
 - Point-to-Point Communication Algorithms.
- The Broadcast Problem
 - Best effort broadcast.
 - Reliable broadcast.

What is a distributed system?

“A distributed system is a network that consists of autonomous computers that are connected using a distribution middleware. They help in sharing different resources and capabilities to provide users with a single and integrated coherent network.” – Technopedia.

What is a distributed system?

“A distributed system is one in which the failure of a computer you didn’t even know existed can render your own computer unusable.”

-- L. Lamport (Computer Science Turing Award)

What is a distributed system?

“A distributed system is one in which the failure of a computer you didn’t even know existed can render your own computer unusable.”

-- L. Lamport (Computer Science Turing Award)

This might be a more honest answer!

What is a distributed system?

“A distributed system is one in which the failure of a computer you didn’t even know existed can render your own computer unusable.”

-- L. Lamport (Computer Science Turing Award)

Note: our goal, as designers of distributed systems and distributed algorithms is to work to ensure that there are definitions that don't depend on failure!


This might be a more honest answer!

What is a distributed system?

A more honest and focused answer:

A distributed system is composed by a set of **processes** that are interconnected through some **network** where processes seek to achieve some form of **cooperation** to execute tasks towards a **common goal**.

What is a distributed system?



Why do we want to distribute things anyways?

Why do we want to distribute things anyways?

- **Fault-Tolerance (also known as *dependability*):** If I use N machines to support my system and f ($f < N$) fail, then my system can still operate.
- **Concurrency (or more processing/storage power):** If instead of using 1 machine to run my system, I use N machines ($N \gg 1$) then I will have N times more resources and, hopefully, my system will be (close to) N times faster.

Challenges of Distributed Systems

There are challenges inherent to all distributed systems:

Challenges of Distributed Systems

There are challenges inherent to all distributed systems:

- **Fault-Tolerance:** It is necessary to consider what you do when a machine fails or becomes unreachable (and they will).

Challenges of Distributed Systems

There are challenges inherent to all distributed systems:

- **Fault-Tolerance:** It is necessary to consider what you do when a machine fails or becomes unreachable (and they will).
- **Concurrency:** It is necessary to consider the possible orderings of events.

Challenges of Distributed Systems

There are challenges inherent to all distributed systems:

- **Fault-Tolerance:** It is necessary to consider what you do when a machine fails or becomes unreachable (and they will).
- **Concurrency:** It is necessary to consider the possible orderings of events.
- **No global clock:** In other words, no way of synchronising via time.

Reasoning about Distributed Systems.

- You can probably guess that distributed systems are complex.

Reasoning about Distributed Systems.

- You can probably guess that distributed systems are complex.
- And our task is to understand how they might behave and how we can build them to operate correctly (and efficiently).

Reasoning about Distributed Systems.

- You can probably guess that distributed systems are complex.
- And our task is to understand how they might behave and how we can build them to operate correctly (and efficiently).
- Uh oh...




The model of a distributed system

- Somewhat “less intelligent” example:
 - How do you reason about a car?

The model of a distributed system

- Somewhat “less intelligent” example:
 - How do you reason about a car?



The model of a distributed system

Key insight:


We need **abstractions** for talking
about all kinds of systems

The model of a distributed system

A distributed system is composed by a set of **processes** that are interconnected through some network (via **links**) where processes seek to achieve some form of cooperation to execute tasks.


The model of a distributed system

A distributed system is composed by a set of **processes** that are interconnected through some network (via **links**) where processes seek to achieve some form of cooperation to execute tasks.



The model of a distributed system


A distributed system is composed by a set of **processes** that are interconnected through some **network (via links)** where processes seek to achieve some form of cooperation to execute tasks.



The model of a distributed system

More formally:

- Processes – computational elements:
 - Abstracts the notion of machine/node.
- Network – Graph $\mathbf{G}=(\mathbf{V},\mathbf{E})$, in which \mathbf{V} is the set of processes, \mathbf{E} represents the communication channels (i.e., links) between pairs of processes.




The model of a distributed system

More formally:


- Processes – computational elements:
 - Abstracts the notion of machine/node.
- Network – Graph $\mathbf{G}=(\mathbf{V},\mathbf{E})$, in which \mathbf{V} is the set of processes, \mathbf{E} represents the communication channels (i.e., links) between pairs of processes.

In general, we will consider a complete graph, where every process is connected to every other by a bidirectional link.




The model of a distributed system

- Processes are fully independent, meaning that they do not share memory in any way.
- But a clear requirement for cooperation is to exchange information... so how do processes exchange information among them?




The model of a distributed system

- Processes are fully independent, meaning that they do not share memory in any way.
- But a clear requirement for cooperation is to exchange information... so how do processes exchange information among them?



The model of a distributed system


- Processes are fully independent, meaning that they do not share memory in any way.
- But a clear requirement for cooperation is to exchange information... so how do processes exchange information among them?



The model of a distributed system

More formally:

- Processes communicate through the exchange of Messages, belonging to an alphabet **M** plus the special symbol **null**, which captures a non-existent message.
- Notation: $\text{send}_i(j, m, \text{arg}_1, \text{arg}_2, \dots, \text{arg}_n)$
 - Process **i** send message **m** with arguments arg_1 to arg_n to process **j**.




The model of a distributed system

- We are done right?

The model of a distributed system

- We are done right?
- Not quite, back to our somewhat “less intelligent” example:
 - How do you reason about a car?






distributed system

ewhat “less intelligent”

How do you



The model of a distributed system

- We are done right?
- Not quite...
- Some extra things must be considered.
 - Timing assumptions.
 - Internal model of the process.
 - (Process) Fault model.
 - Network model,

Examples

- Client-server web interactions
- Audio streaming
- Stock-exchange trading

Timing assumptions

Two fundamental models

Synchronous System:

- Assumes that there is a known *upper bound* to the time required to deliver a message through the network and for a process to make all computations related with the processing of the message.

Asynchronous System:

- There are *no assumptions* about the time required to deliver a message or process a message.

Timing assumptions

- This might look like not a big deal but, there are strong implications here:
 - In a synchronous system you can **detect** when a process fails (in some particular fault models).
 - In a synchronous system you can have protocols **evolve** in synchronous steps (Why is that?)
 - In an asynchronous system there are some problems that have **no solution**.

Timing assumptions

The “real world” is **asynchronous**, so why is it that we sometimes consider the synchronous model?

- The synchronous model is **easier** to reason about
- More **efficient** solutions can be devised in the synchronous model

Internal Model of the Process

- Each process has a unique identifier: $i \in V$
- Internally each process has (classical model):
 - $states_i$ – set of (valid) states for process i
 - Inputs and outputs that are special state variables (which allow the process to get information from outside and export information to the outside)
 - $init_i$ – initial state for process i
 - Fundamentally, a process is a **deterministic automaton**.

Beyond Synchrony and Asynchrony

- There is a (more) practical system model that in some sense stands between these two extremes.
- The **partially synchronous** model (or *eventually synchronous* model).
- In this model the system is considered to be **asynchronous**, but it is assumed that **eventually** (*meaning for sure at some time in the future - that is unknown*) the system will behave in a **synchronous** way for long enough (for something good to happen).

Transition between States

- Synchronous Model: Execution in rounds.
- Transition of the process (internal) state based on two functions:
 - $\text{trans}_i : \text{states}_i \times [M \cup \{\text{null}\}]^V \rightarrow \text{states}_i$
 - $\text{msgs}_i : \text{states}_i \rightarrow [M \cup \{\text{null}\}]^V$
- In each round a process will:
 - Sends a message (potentially different) to all processes.
 - Receives a message from all processes.
 - Apply trans_i over the received messages to determine its following state.

Transition between States

- Synchronous Model: Execution in rounds.
- Transition of the process (internal) state based on two functions:
 - $\text{trans}_i : \text{states}_i \times [M \cup \{\text{null}\}]^V \rightarrow \text{states}_i$
 - $\text{msgs}_i : \text{states}_i \rightarrow [M \cup \{\text{null}\}]^V$
- In each round a process will:
 - **Sends a message (potentially different) to all processes.**
 - Receives a message from all processes.
 - Apply trans_i over the received messages to determine its following state.

Transition between States

- Synchronous Model: Execution in rounds.
- Transition of the process (internal) state based on two functions:
 - $\text{trans}_i : \text{states}_i \times [\mathbf{M} \cup \{\mathbf{null}\}]^V \rightarrow \text{states}_i$
 - $\text{msgs}_i : \text{states}_i \rightarrow [\mathbf{M} \cup \{\mathbf{null}\}]^V$
- In each round a process will:
 - Sends a message (potentially different) to all processes.
 - **Receives a message from all processes.**
 - Apply trans_i over the received messages to determine its following state.


Transition between States

- Synchronous Model: Execution in rounds.
- Transition of the process (internal) state based on two functions:
 - $\text{trans}_i : \text{states}_i \times [M \cup \{\text{null}\}]^V \rightarrow \text{states}_i$
 - $\text{msgs}_i : \text{states}_i \rightarrow [M \cup \{\text{null}\}]^V$
- In each round a process will:
 - Sends a message (potentially different) to all processes.
 - Receives a message from all processes.
 - **Apply trans_i over the received messages to determine its following state.**


Transition between States

- Asynchronous Model: Execution is not based on rounds.
- Transition of the process state based on two functions:
 - $\text{trans}_i : \text{states}_i \times M \cup \{\text{null}\} \rightarrow \text{states}_i$
 - $\text{msgs}_i : \text{states}_i \rightarrow \{M \cup \{\text{null}\}\}^V$
- Since there is no notion of rounds:
 - A transition of state is triggered by the reception of a single message (notice that the transition can be $S \rightarrow S$)
 - Transitioning to a state (even if the same) can trigger the generation (and transmission) of a new set of messages.

Example of an execution in the synchronous model:




Example of an execution in the asynchronous model:



Fault Model

- Faults lead processes to deviate from their expected behavior.
- Classical Model:



- Example: Sector in the hard disk is damaged (fault);
-> Sector is accessed (Error) -> File is lost (Failure)

Fault Model (cont.)

- This classical model usually has a recursive implication.
- The failure of a component of a process (or system) might imply a fault in another component (or different process in the system).
- Going back to the previous example, the failure of the file system (file damaged) might lead to a fault in the load of the operative system, which might result in the failure of the operative system.

Process Fault Model

- A process that never fails, is considered **correct**.
- **Correct processes** never deviate from their expected/prescribed behavior.
 - It executes the algorithm as expected and sends all messages defined by it.
- **Failed (or Faulty) processes** might deviate from their prescribed behavior in different ways.
 - The *unit of failure* is the process, meaning that when it fails, all its component fail at the same time.
- The (possible/considered) behaviors of a faulty process is defined by the process fault model.

Process Fault Models

- Crash Fault Model:
 - When a process fails it stops sending any messages (from that point onward).
 - This is the fault model that we will consider most of the times.

Process Fault Models

- Crash Fault Model:
 - When a process fails it stops sending any messages (from that point onward).
 - This is the fault model that we will consider most of the times.
- Omission Fault Model:
 - A process that fails omits the transmission (or reception) of any number of messages (potentially not all of them).

Process Fault Models

- Crash Fault Model:
 - When a process fails it stops sending any messages (from that point onward).
 - This is the fault model that we will consider most of the times.
- Omission Fault Model:
 - A process that fails omits the transmission (or reception) of any number of messages (potentially not all of them).
- Fail-Stop Model:
 - Similar to the crash model, except that upon failure the process “notifies” all other processes of its own failure (only possible when considering a synchronous system).

Process Fault Models

- Byzantine (or Arbitrary) Fault Model:
 - A failed process might deviate from its protocol in any arbitrary way.
 - Examples:
 - Duplicate Messages;
 - Create invalid messages;
 - Modify values received from other processes.
 - Why is this relevant?
 - Can capture memory corruption;
 - Can capture software bugs;
 - A malicious attacker that controls a process.

Process Fault Models

- Byzantine (or Arbitrary) Fault Model:
 - A failed process might deviate from its protocol in any arbitrary way.
 - Examples:
 - Duplicate Messages;
 - Create invalid messages;
 - Modify values received from other processes.
 - Why is this relevant?

Process Fault Models

- Byzantine (or Arbitrary) Fault Model:
 - A failed process might deviate from its protocol in any arbitrary way.
 - Examples:
 - Duplicate Messages;
 - Create invalid messages;
 - Modify values received from other processes.
 - Why is this relevant?
 - Can capture memory corruption;
 - Can capture software bugs;
 - A malicious attacker that controls a process.
 - This is not something that we will delve into here:
 - The **Confiabilidade de Sistemas Distribuídos (CSD)** course deals with these challenges.

Network Model

- The Network Model captures the assumptions made concerning the **links** that interconnect processes.
- Namely it captures what can go wrong in the network regarding:
 - Messages sent between processes being lost.
 - Possibility of duplication of messages.
 - Possibility for corruption of messages.

Network Model

- The Network Model captures the assumptions made concerning the **links** that interconnect processes.
- Namely it captures what can go wrong in the network regarding:
 - Messages sent between processes being lost.
 - Possibility of duplication of messages.
 - Possibility for corruption of messages.

Network Model

- The Network Model was made conceptually precise

The Network is usually not your friend.

Network Model

- The Network Model was made conceptually precise

The Network is usually not your friend.

Some of the weird phenomena in the network can be understood in the **Architectures and Protocols of Computer Networks (APRC)** course.

How to model the network?

We need abstractions like we had with the timing and fault models

In principle, we make a set of assumptions about the network operates

Some models ***closer*** to reality than others

Network Model (a starting point)

Fair Loss Model

- A model that captures the possibility of messages **being lost** (albeit in a fair way)
- Properties:
 - FL1 (Fair-Loss): Considering two correct processes i and j ; if i sends a message to j infinite times, then j delivers the message infinite times.
 - FL2 (Finite Duplication): Considering two correct processes i and j ; if i sends a message m to j a finite number of times, then j cannot deliver m infinite times.
 - FL3 (No Creation): If a correct process j delivers a message m , then m was sent to j by some process i .

Network Model (moving on...)

Stubborn Model

- A stronger model that assumes that processes communicate in a stubborn way, to **prevent message loss**.
- Properties:
 - SL1 (Stubborn Delivery): Considering two correct processes i and j ; if i sends a message to j , then j delivers the message an infinite number of times.
 - SL2 (No Creation): If a correct process j delivers a message m , then m was sent to j by some process i .

Network Model (better now)

Perfect Link Model (also called Reliable):

- A stronger model that assumes the links between processes are well-behaved.
- Properties:
 - PL1 (Reliable Delivery): Considering two correct processes i and j ; if i sends a message to j , then j ***eventually*** delivers m .
 - PL2 (No Duplication): No message is delivered by a process more than once.
 - PL3 (No Creation): If a correct process j delivers a message m , then m was sent to j by some process i .

How does all this fits in?

- Our typical network is more closely captured by the properties of the fair-loss model, however its frequent that we use the perfect link model... Why?

How does all this fits in?

- Our typical network is more closely captured by the properties of the fair-loss model, however its frequent that we use the perfect link model... Why?
 - The perfect link model makes it easier to reason about algorithms design...

How does all this fits in?

- Our typical network is more closely captured by the properties of the fair-loss model, however its frequent that we use the perfect link model... Why?
 - The perfect link model makes it easier to reason about algorithms design...
 - ...but more importantly, these abstractions can be built on top of one another using **distributed algorithms**.

How does all this fits in?

- Our proposed frequency model is more closely captured by the loss model, however its the perfect link model... Why?
 - The model makes it easier to reason about
 - ., these abstractions can be built using **distributed algorithms**.
- **What is this dark sorcery that you speak?**



How does all this fits in?

- Our
pro
frec



-
-

more closely captured by the
oss model, however its
he perfect link model... Why?

l makes it easier to reason about

y, these abstractions can be built
using **distributed algorithms**.

- **What is this dark sorcery that you speak?**
- **Wait for the lab today to discover...**

How does all

Don't forget that the perfect link abstraction does not state anything about time

Algorithms Specification and Properties

- Did you noticed that when discussing these network models (i.e., abstractions) I have defined them in terms of a set of properties?

Algorithms Specification and Properties

- Did you noticed that when discussing these network models (i.e., abstractions) I have defined them in terms of a set of properties?
- **Algorithms** (that instantiate these abstractions) also provide a set of **properties** (if correct, those of the abstraction they implement).
- Why do we tend to think in terms of properties?

Algorithms Specification and Properties

- Did you noticed that when discussing these network models (i.e., abstractions) I have defined them in terms of a set of properties?
- **Algorithms** (that instantiate these abstractions) also provide a set of **properties** (if correct, those of the abstraction they implement).
- Why do we tend to think in terms of properties?
- Quick answer: Because algorithms are ***composable***, and the design of an algorithm depends on the underlying properties provided by other algorithms.

Algorithms Specification and Properties

What does these properties capture?

- The correctness criteria for the algorithm (and for any implementation of that algorithm).
- It defines restrictions to (all) valid executions of the algorithm.

Two fundamental types of properties:

- Safety
- Liveness

Safety Properties

Conditions that must be enforced at any (and all) point of the execution

- Intuitively, preventing *bad* things that should never happen.

Relevant aspects:

- The trace of an empty execution is always safe (do nothing and you shall do nothing wrong).
- The prefix of a trace that does not violate safety, will never violate safety.

Liveness Properties

Conditions that should be enforced at some point of an execution (but not necessarily always).

- Intuitively, good things that should happen eventually.

Relevant aspects:

- One can always extend the trace of an execution in a way that will respect liveness conditions (if you haven't done anything good yet, you might do it in the future).

Safety VS Liveness Properties

- Correct algorithms will have (most of the times) both Safety and Liveness properties.
- Some properties however are hard to classify within one of these classes, and they might mix aspects of safety and liveness
 - Usually, one can decompose these properties in simpler ones through conjunctions.

Modelling distributed systems

Models that handle:

- Networks
- Timing
- Faults

Currently, we do not have any constructions that actually **operate** within these models

Let's now start to think about concrete problems...

- Now that we have covered the rules of the game, we can start to think on how to solve problems (within a particular set of rules).

The Broadcast Problem

- Informally: A process needs to transmit the same message m to N other processes (where N is every process in the system including himself).
- *Let's assume that the complete set of processes in the system is known a-priori: π*
- *Let's assume we have access to the Perfect Point-to-Point Link Abstraction.*
- *Let's assume an asynchronous system (no rounds, no failure detection).*

The Broadcast Problem

- Wait... How do you specify an algorithm again?
- I have sort of told you before, through a deterministic automaton...

The Broadcast Problem

- Wait... How do you specify an algorithm again?
- I have sort of told you before, through a deterministic automaton...
- Ok let's simplify this...

The Anatomy of an Algorithm:



ALGORITHM

The Anatomy of an Algorithm:



ALGORITHM

You need an Interface for this to be used (think of APIs)

The Anatomy of an Algorithm:



ALGORITHM

You need an Interface for this to
be used (think of APIs)

Let's use an event driven interface.

The Anatomy of an Algorithm:

Perfect Point-to-Point Link Algorithm

You need an Interface for this to
be used (think of APIs)

Let's use an event driven interface.

The Anatomy of an Algorithm:

pp2pSend (p , m)

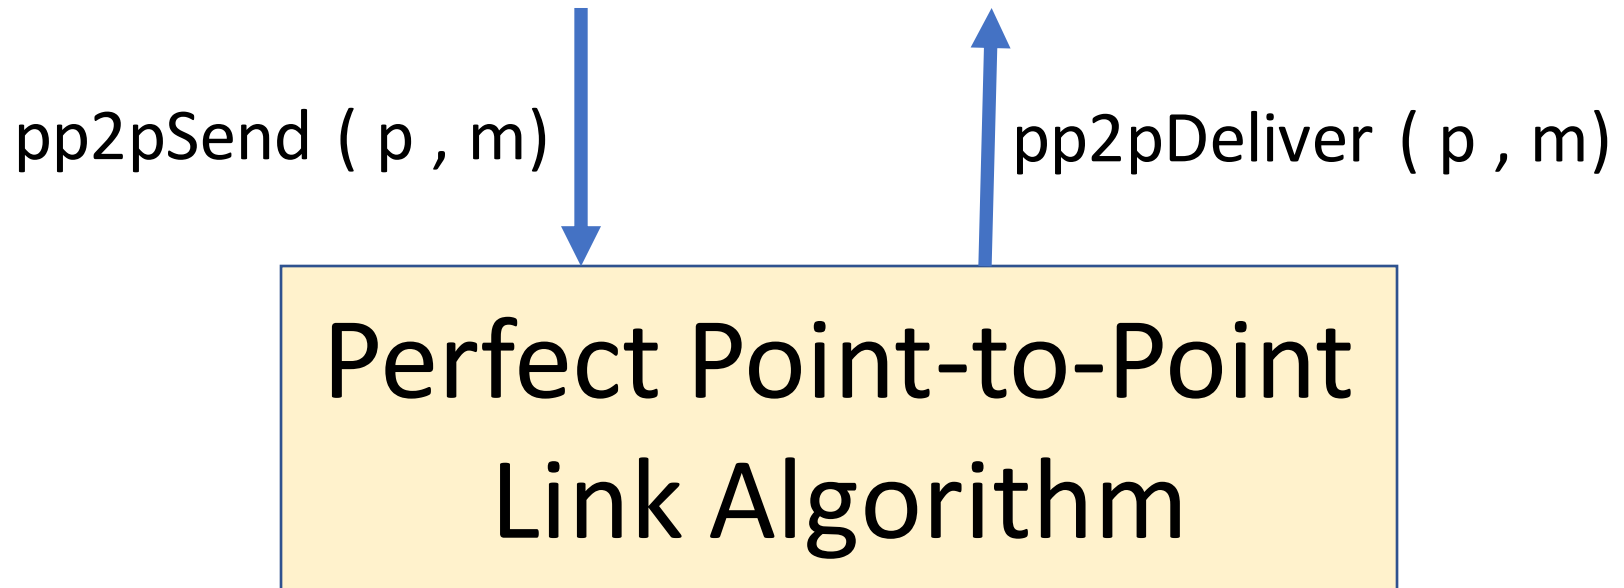


**Perfect Point-to-Point
Link Algorithm**

You need an Interface for this to
be used (think of APIs)

Let's use an event driven interface.

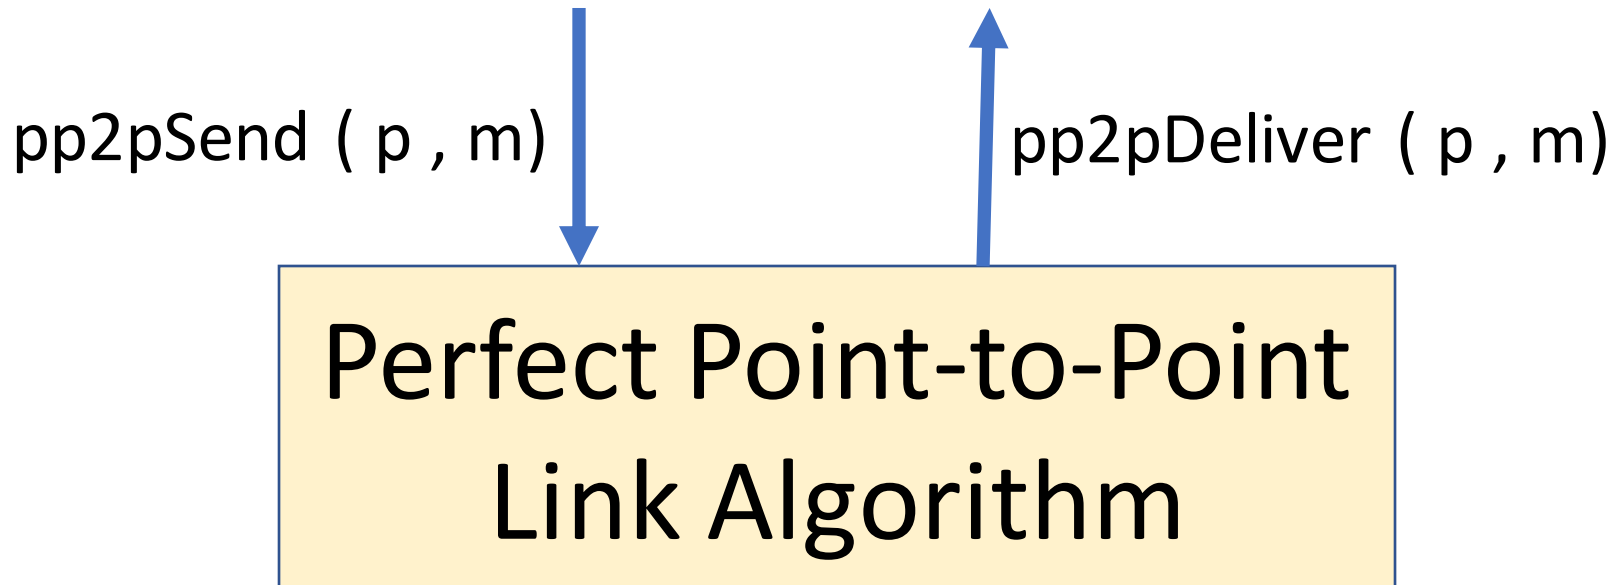
The Anatomy of an Algorithm:



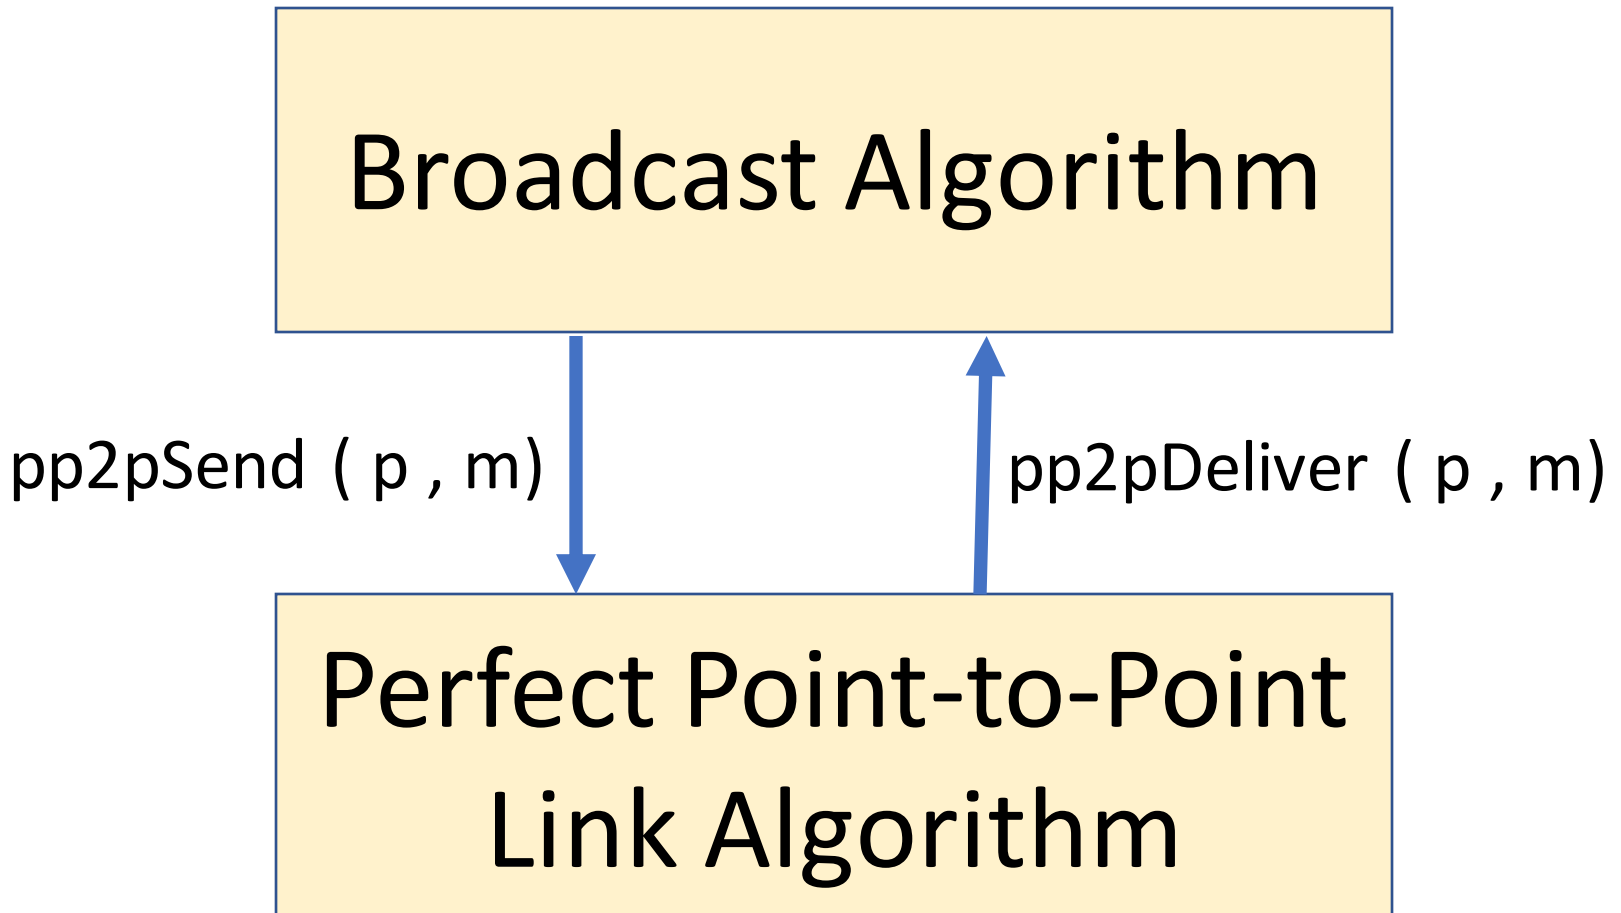
You need an Interface for this to be used (think of APIs)

Let's use an event driven interface.

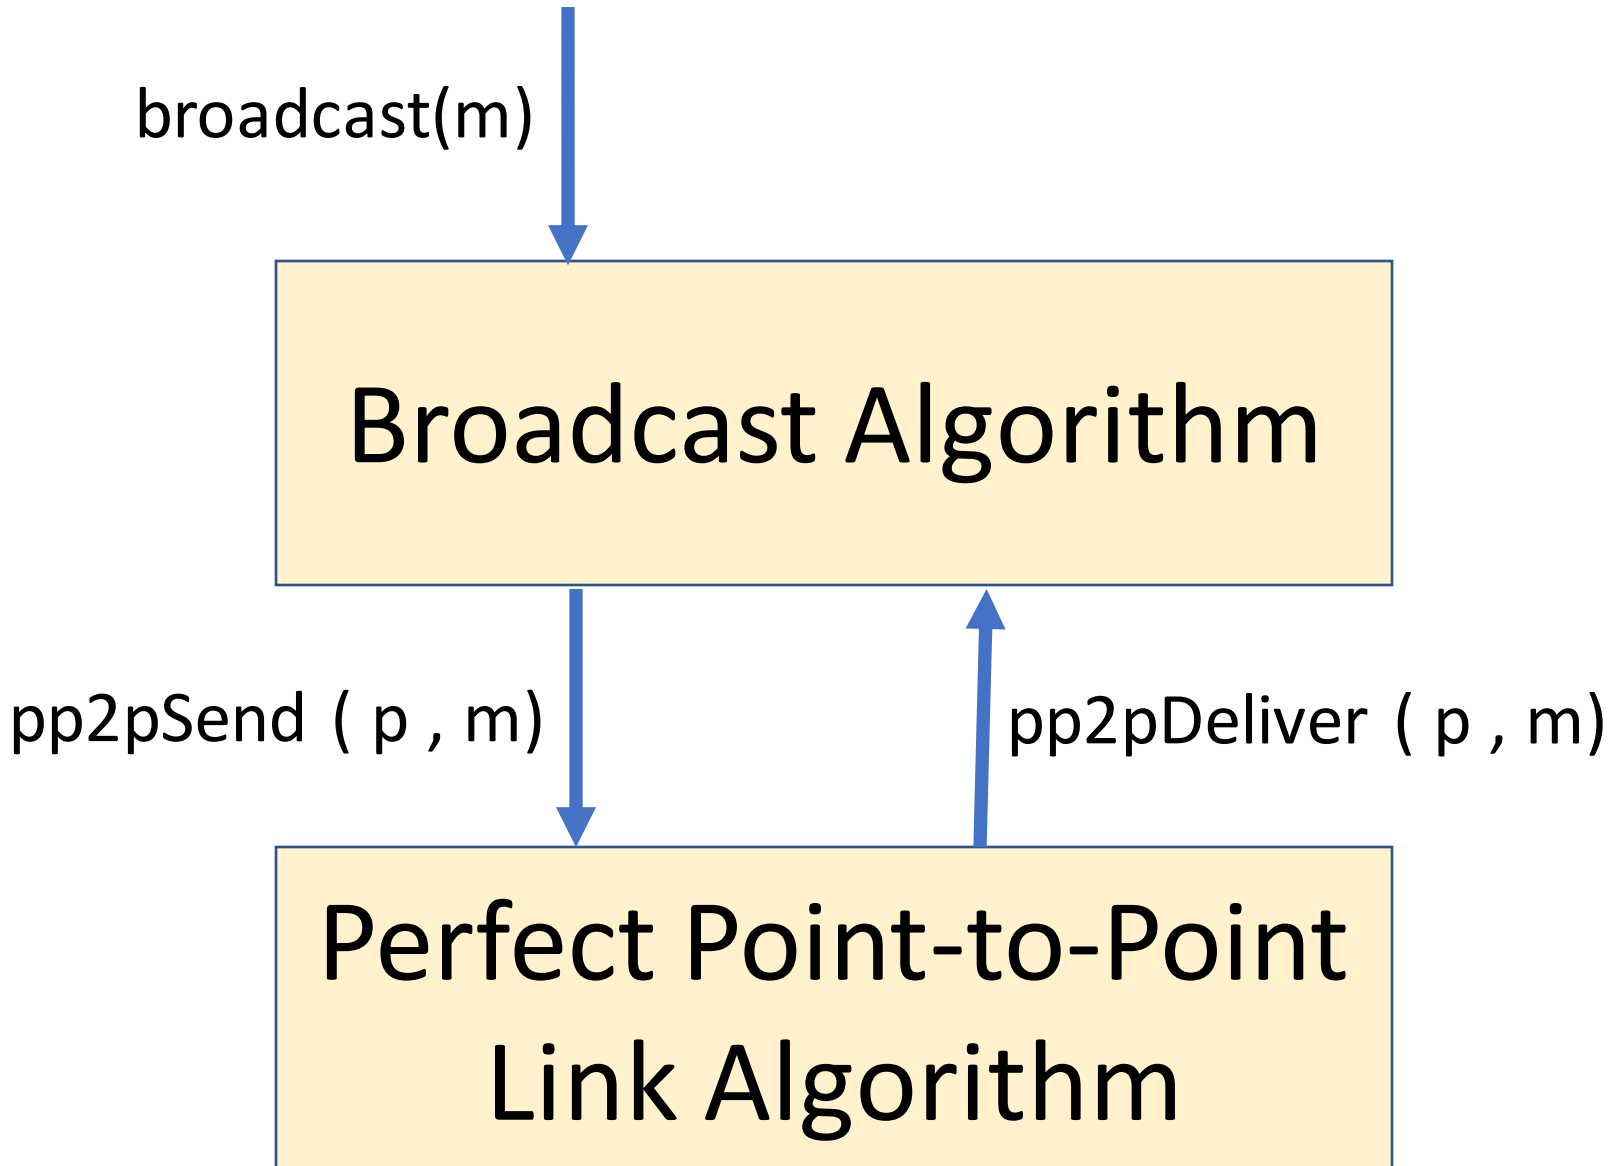
The Anatomy of an Algorithm:



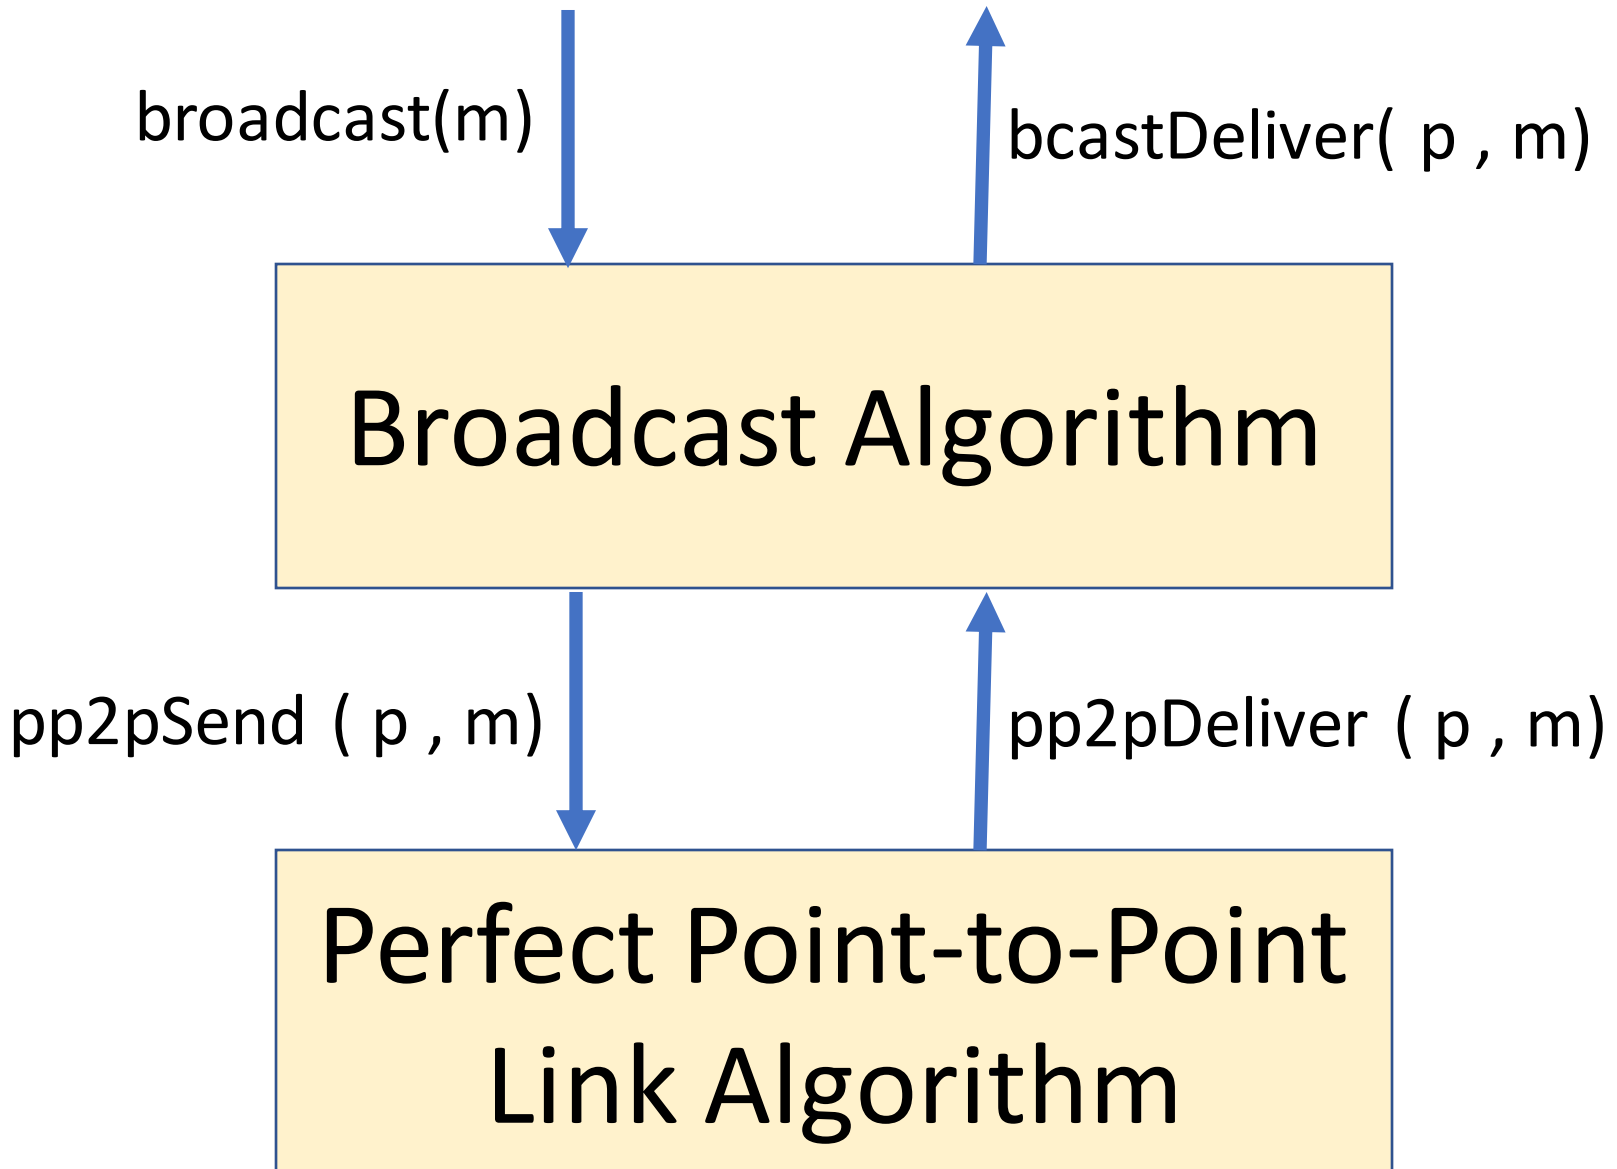
The Anatomy of an Algorithm:



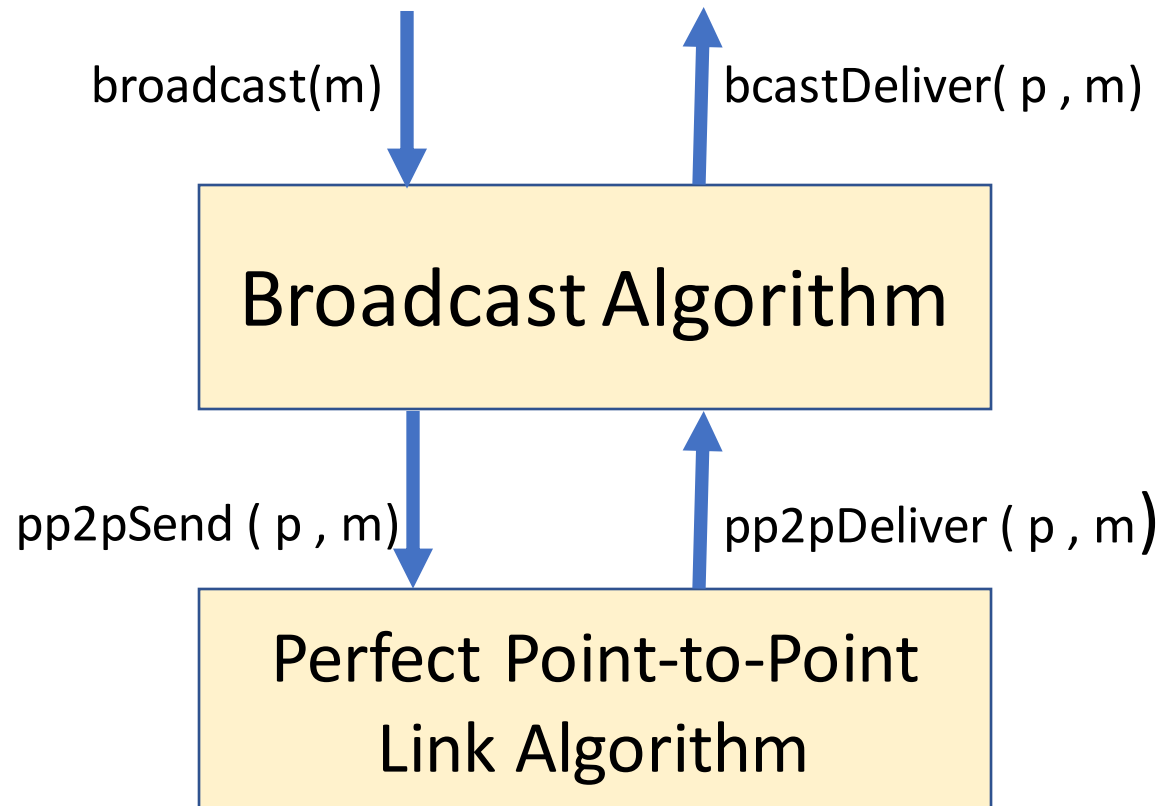
The Anatomy of an Algorithm:



The Anatomy of an Algorithm:



The Anatomy of an Algorithm:




Now you specify what actions are performed when one of these events happen:

Upon `broadcast(m)` do:

...

...

The Anatomy of an Algorithm:




Now you specify what actions are performed when one of these events happen:

Upon broadcast(m) do:
...
...

You can trigger an event on other protocol:
Trigger pp2psend(p, m)

The Anatomy of an Algorithm:

There is a special event called **Init** where you can initialize the **local** state of the **algorithm**



Now you specify what happens when one of these events happen:

Upon broadcast(m) do:

...


You can trigger an event on another protocol:

Trigger pp2psend(p, m);

Layering (Broadcasts)


Two important distinctions:

- Requests
- Indications




Requests

Requests inform the *next* layered algorithm to perform some action.



Indications

Indications inform the *previous* layered algorithm that the action was performed.



Back to the Broadcast problem...

- What is the simplest solution that you can think of?
- **Think like this:**
 - I am a process; someone asks me to send a message to everyone (including me)
 - I know everyone already so basically what do I do?

Back to the Broadcast problem...

- What is the simplest solution that you can think of?

- **Solution:**

Just go ahead and send the message to everyone, one at a time. When you get one of these messages, you just deliver it to the upper layer.

Back to the Broadcast Problem:

Good... that works... (Sort of)

- That is the solution for the Best Effort Broadcast.
- Best Effort Broadcast:
 - BEB1 (Best-Effort validity): For any two correct processes i and j , every message broadcasted by i is eventually delivered by j .
 - BEB2: (No Duplication): No message is delivered more than once.
 - BEB3: (No Creation): If a correct process j delivers a message m , then m was **broadcast** by some process i .

Best Effort Broadcast

State: (could be omitted)

Upon Init do: (could be omitted)

Upon bebBroadcast(m) **do**
 forall $p \in \pi$
 trigger pp2pSend(p, m);

Upon pp2pDeliver (p, m) **do**
 trigger bebDeliver(p, m);

Best Effort Broadcast

- Not so great right?
- What happens if a process fails while sending messages to all the other processes?

Best Effort Broadcast

- Not so great right?
- What happens if a process fails while sending messages to all the other processes?
- Answer not everyone gets the message...

Reliable Broadcast Problem

- Let's make this somewhat stronger and more interesting...
- Reliable Broadcast:
 - RB1 (Validity): If a correct process i broadcasts message m , then i eventually delivers the message.
 - RB2 (No Duplications): No message is delivered more than once.
 - RB3 (No Creation): If a correct process j delivers a message m , then m was broadcast to j by some process i .
 - RB4 (Agreement): If a message m is delivered by some correct process i , then m is eventually delivered by every correct process j .

How do you think we could solve this problem?

Remember that we can use the Best Effort Broadcast Algorithm to solve this one...

- Let's make this somewhat stronger and more interesting...
- **Reliable Broadcast:**
 - RB1 (Validity): If a correct process i broadcasts message m , then i eventually delivers the message.
 - RB2 (No Duplications): No message is delivered more than once.
 - RB3 (No Creation): If a correct process j delivers a message m , then m was broadcast to j by some process i .
 - RB4 (Agreement): If a message m is delivered by some correct process i , then m is eventually delivered by every correct process j .

Reliable Broadcast Problem

State:

delivered //set of message ids that were already delivered.

Upon Init do:

delivered $\leftarrow \{\}$;

Upon rbBroadcast(m) **do**

trigger rbDeliver(m);

mid \leftarrow generateUniqueID(m);

delivered \leftarrow delivered \cup {mid};

trigger bebBroadcast({ mid, m });

Upon bebDeliver(p, { mid, m }) **do**

if (mid \notin delivered) **then**

delivered \leftarrow delivered \cup {mid};

trigger rbDeliver(m);

trigger bebBroadcast(mid, m);

Reliable Broadcast Problem

In all fairness in the Literature: *Eager Reliable Broadcast*

State:

delivered //set of message ids that were already delivered.

Upon Init **do**:

delivered $\leftarrow \{\}$;

Upon rbBroadcast(m) **do**

trigger rbDeliver(m);

mid \leftarrow generateUniqueID(m);

delivered \leftarrow delivered \cup {mid};

trigger bebBroadcast({ mid, m });

Upon bebDeliver(p, { mid, m }) **do**

if (mid \notin delivered) **then**

delivered \leftarrow delivered \cup {mid};

trigger rbDeliver(m);

trigger bebBroadcast(mid, m);

Reliable Broadcast Problem

In all fairness in the Literature: *Eager Reliable Broadcast*

State:

delivered //set of message ids that were already delivered.

Upon Init do:

delivered $\leftarrow \{\}$;

Upon rbBroadcast(m) do

trigger rbDeliver(m);

mid \leftarrow generateUniqueID(m);

delivered \leftarrow delivered \cup {mid};

trigger bebBroadcast({ mid, m });

Upon bebDeliver(p, { mid, m }) do

if (mid \notin delivered) **then**

delivered \leftarrow delivered \cup {mid};

trigger rbDeliver(m);

trigger bebBroadcast(mid, m);

Why is this solution correct?

To show the correctness of the algorithm we must check each property individually. For each we have to make a (rational) argument that the algorithm enforces it.

- **Reliable Broadcast:**
 - RB1 (Validity): If a correct process i broadcasts message m , then i eventually delivers the message.
 - RB2 (No Duplications): No message is delivered more than once.
 - RB3 (No Creation): If a correct process j delivers a message m , then m was broadcast to j by some process i .
 - RB4 (Agreement): If a message m is delivered by some correct process i , then m is eventually delivered by every correct process j .

Reliable Broadcast Problem

In all fairness in the Literature: *Eager Reliable Broadcast*

State:

delivered //set of message ids that were already delivered.

Upon Init do:

delivered $\leftarrow \{\}$;

Upon rbBroadcast(m) do

trigger rbDeliver(m);

mid \leftarrow generateUniqueID(m);

delivered \leftarrow delivered \cup {mid};

trigger bebBroadcast({ mid, m });

Upon bebDeliver(p, { mid, m }) do

if (mid \notin delivered) **then**

delivered \leftarrow delivered \cup {mid};

trigger rbDeliver(m);

trigger bebBroadcast(mid, m);

This solution consumes a lot of messages (not being nice for the network nor processors)!

In fact, it costs $O(|\pi|^2)$ total messages to execute.

Reliable Broadcast Problem

- Ideally we would design a Reliable Broadcast solution that resulted in $O(|\pi|)$ messages
- See Lab 01!

Summary of Lecture

- Why are distributed systems required?
- The importance of abstractions
- Idealised models for describing distributed algorithms
- Concrete examples of broadcasting messages

Quiz!!1!



What do you remember?



Question 1

Which p2p link model allows for the loss of messages?

- Fair loss
- Stubborn
- Perfect

Answer 1

Which p2p link model allows for the loss of messages?

- **Fair loss**
- Stubborn
- Perfect

Question 2

In the synchronous model, it is impossible to tell when a process has failed?

- True
- False

Answer 2

In the synchronous model, it is impossible to tell when a process has failed?

- True
- **False**

Question 3

Which fault model allows a process to keep sending and delivering a subset of the messages it receives?

- Crash fault
- Omission fault
- Fail-stop

Answer 3

Which fault model allows a process to keep sending and delivering a subset of the messages it receives?

- Crash fault
- **Omission fault**
- Fail-stop